

# PROBLEM STATEMENT

For a health insurance company to make money, it needs to collect more in yearly premiums than it spends on medical care to its beneficiaries. As a result, insurers invest a great deal of time and money in developing models that accurately forecast medical expenses for the insured population.

The goal of this analysis is to use patient data to estimate the average medical care expenses for such population segments. These estimates can be used to create actuarial tables that set the price of yearly premiums higher or lower, depending on the expected treatment costs.

## IMPORTING LIBRARIES

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

In [2]:

```
insurance_data = pd.read_csv(r"C:\Users\shaha\OneDrive\Desktop\Excel\insurance.csv")
insurance_data
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [3]:

```
insurance_data.head()
```

Out[3]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [4]:

```
insurance_data.tail()
```

Out[4]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [5]:

```
insurance_data.shape
```

Out[5]:

(1338, 7)

In [6]:

```
insurance_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [7]:

```
insurance_data.isnull().sum()
```

Out[7]:

```
age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```

In [8]:

```
insurance_data.describe()
```

Out[8]:

	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010

In [9]:

```
insurance_data.columns
```

Out[9]:

```
Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype  
='object')
```

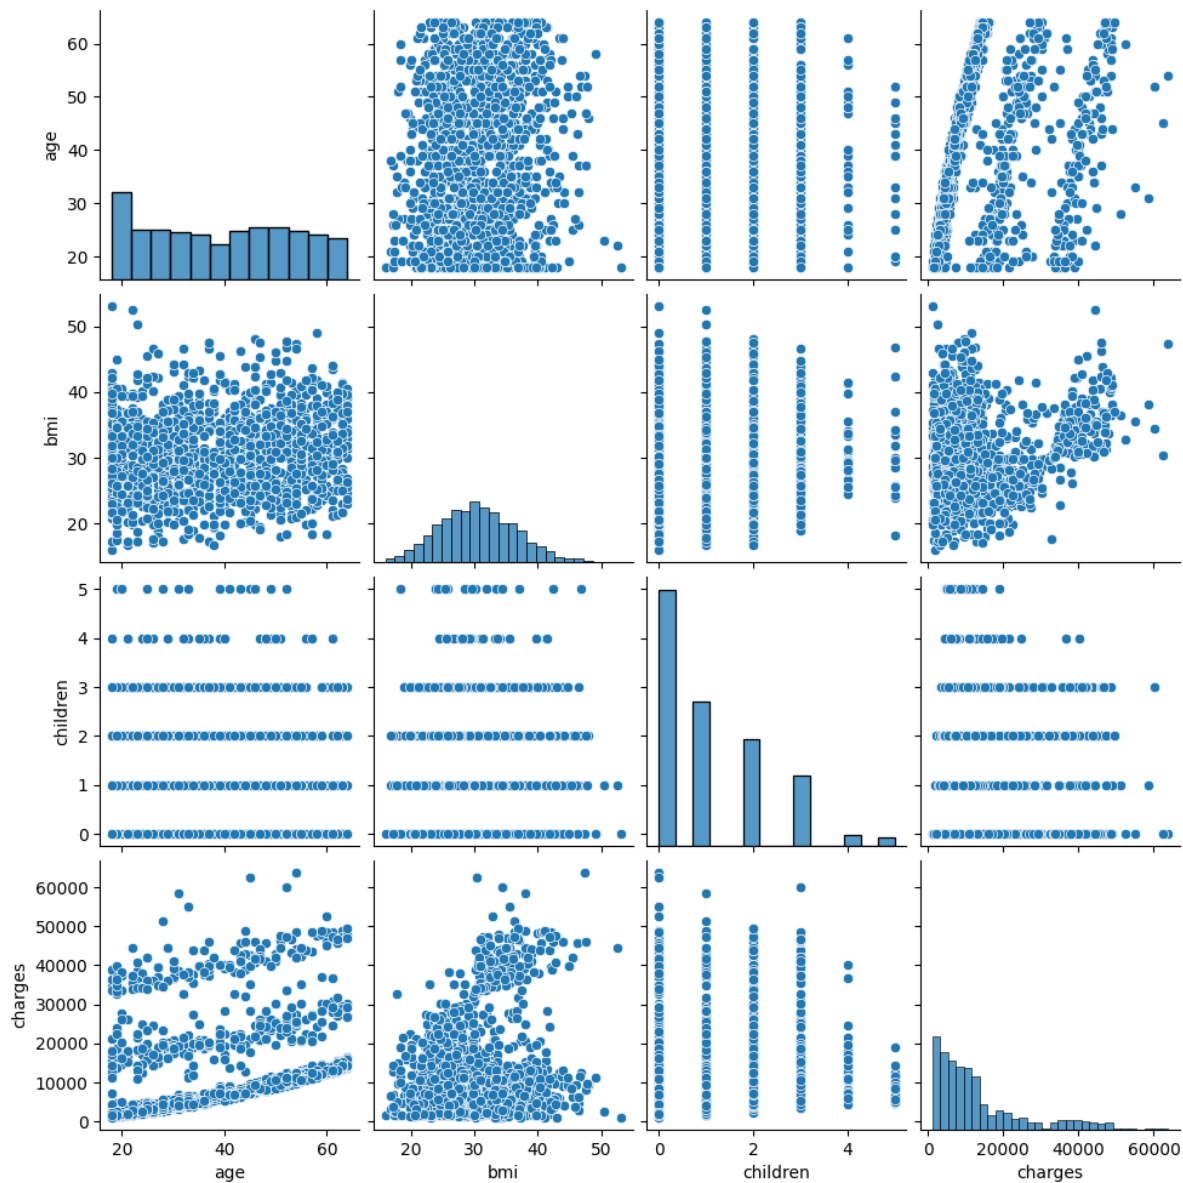
## Exploratory Data Analysis(Linear Regression)

In [10]:

```
sns.pairplot(insurance_data)
```

Out[10]:

<seaborn.axisgrid.PairGrid at 0x25c33ca7b50>




In [11]:

```
sns.displot(insurance_data['age'])
```

Out[11]:

<seaborn.axisgrid.FacetGrid at 0x25c67753550>




In [12]:

```
sns.displot(insurance_data['charges'])
```

Out[12]:

<seaborn.axisgrid.FacetGrid at 0x25c33d5f520>



In [13]:

```
s={"sex":{"female":1,"male":0}}
insurance_data=insurance_data.replace(s)
insurance_data
```

Out[13]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	0	33.770	1	no	southeast	1725.55230
2	28	0	33.000	3	no	southeast	4449.46200
3	33	0	22.705	0	no	northwest	21984.47061
4	32	0	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [14]:

```
som={"smoker":{"yes":1,"no":0}}
insurance_data=insurance_data.replace(som)
insurance_data
```

Out[14]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	0	33.770	1	0	southeast	1725.55230
2	28	0	33.000	3	0	southeast	4449.46200
3	33	0	22.705	0	0	northwest	21984.47061
4	32	0	28.880	0	0	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

In [15]:

```
insurance_data['region'].value_counts()
```

Out[15]:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [16]:

```
reg={"region":{"southeast":1,"southwest":2,"northwest":3,"northeast":4}}
insurance_data=insurance_data.replace(reg)
insurance_data
```

Out[16]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	2	16884.92400
1	18	0	33.770	1	0	1	1725.55230
2	28	0	33.000	3	0	1	4449.46200
3	33	0	22.705	0	0	3	21984.47061
4	32	0	28.880	0	0	3	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	0	3	10600.54830
1334	18	1	31.920	0	0	4	2205.98080
1335	18	1	36.850	0	0	1	1629.83350
1336	21	1	25.800	0	0	2	2007.94500
1337	61	1	29.070	0	1	3	29141.36030

1338 rows × 7 columns

In [17]:

```
df=insurance_data[['age','sex','bmi','children','smoker','region','charges']]
```

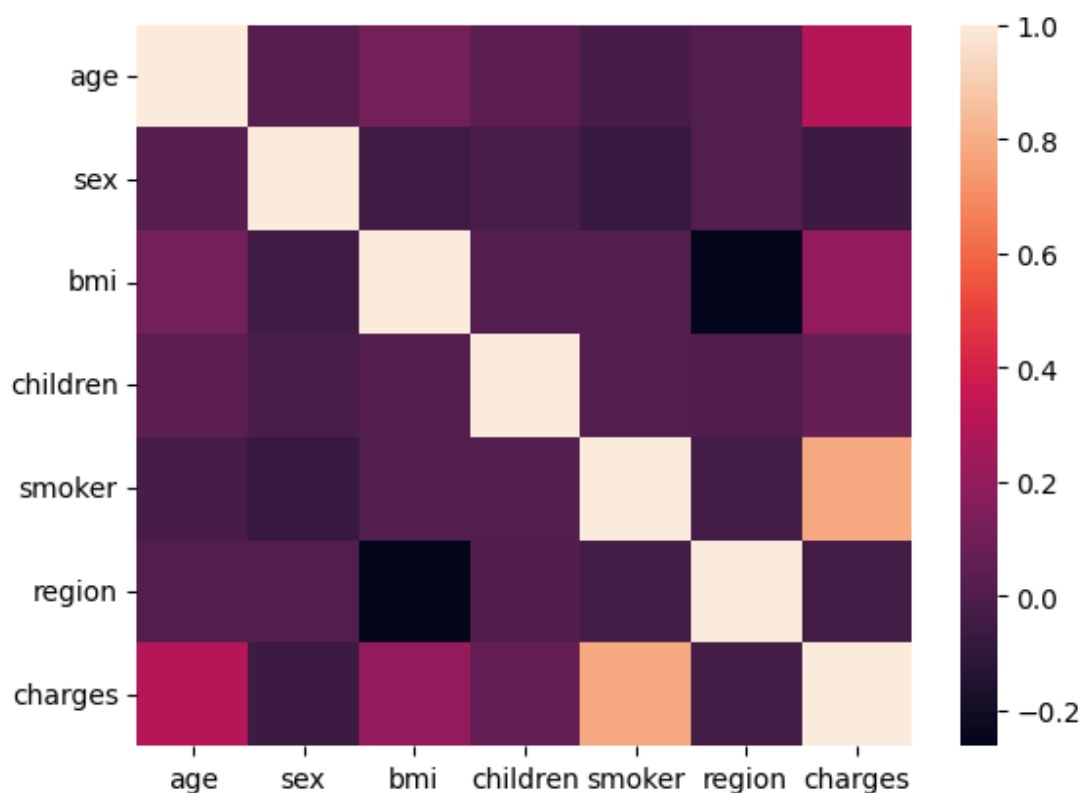


In [18]:

```
sns.heatmap(df.corr())
```

Out[18]:

&lt;Axes: &gt;



## To Train The model

we are going to train linear Regression model. we need to first split up our data into x list that contains the features to train on, and y list with the target variable.

In [19]:

```
x=df[['age','sex','bmi','children','smoker','region']]
y=df['charges']
```

In [20]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=101)
```

In [21]:

```
lm=LinearRegression()
lm.fit(x_train,y_train)
```

Out[21]:

```
LinearRegression()
LinearRegression()
```

In [22]:

```
print(lm.intercept_)
```

-13563.198964426998

In [23]:

```
coeff_df=pd.DataFrame(lm.coef_,x.columns,columns=['coefficient'])  
coeff_df
```

Out[23]:


	coefficient
<b>age</b>	240.808027
<b>sex</b>	-57.898004
<b>bmi</b>	369.026700
<b>children</b>	494.438485
<b>smoker</b>	23470.199126
<b>region</b>	281.762820

In [24]:

```
predictions=lm.predict(x_test)  
plt.scatter(y_test,predictions)
```

Out[24]:

<matplotlib.collections.PathCollection at 0x25c69494550>




In [25]:

```
sns.displot((y_test-predictions),bins=50)
```

Out[25]:

<seaborn.axisgrid.FacetGrid at 0x25c694bc790>



In [26]:

```
from sklearn.metrics import r2_score
y_pred=lm.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.7621231693371181

## Logistic Regression

In [27]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

In [28]:

```
pd.set_option('display.max_rows',1000000000)
pd.set_option('display.max_columns',1000000000)
pd.set_option('display.width',95)
```

In [29]:

```
print('This DataFrame has %d Rows and %d columns'%(df.shape))
```

This DataFrame has 1338 Rows and 7 columns

In [30]:

```
features_matrix=df.iloc[:,0:3]
```

In [31]:

```
target_vector=df.iloc[:,-3]
```

In [32]:

```
print('The Features Matrix Has %d Rows And %d Column(s)'%(features_matrix.shape))
print('The Target Matrix Has %d Rows and %d columns(s)'%(np.array(target_vector).reshape(-1,1)
```

The Features Matrix Has 1338 Rows And 3 Column(s)

The Target Matrix Has 1338 Rows and 1 columns(s)

In [33]:

```
features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

In [34]:

```
algorithm=LogisticRegression(max_iter=1000000)
```

In [35]:

```
logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

In [36]:

```
observation=[[1,0,0.3]]
```

In [37]:

```
predictions=logistic_Regression_Model.predict(observation)
print('The Model predicted the observation to belong to class %s'%(predictions))
```

The Model predicted the observation to belong to class [0]

In [38]:

```
print('The algorithm was trained to predict one of the two classes:%s'%(algorithm.classes_))
```

The algorithm was trained to predict one of the two classes:[0 1]

In [39]:

```
print("""The model says the probability of the obserbvation we passedbelonging to class['b']is
      %(algorithm.predict_proba(observation)[0][0]))
print()
print("""The model says the probability of the observation we passed belonging to class['g']is
      %(algorithm.predict_proba(observation)[observation[0][1]]))
```

The model says the probability of the obserbvation we passedbelonging to class  
['b']is 0.806147490103992

The model says the probability of the observation we passed belonging to class  
['g']is [0.80614749 0.19385251]

In [40]:

```
x=np.array(df['age']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-3,1)
```

In [41]:

```
lr=LogisticRegression()
lr.fit(x,y)
print(lr.score(x,y))
```

0.7952167414050823

C:\Users\shaha\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn  
n\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed  
when a 1d array was expected. Please change the shape of y to (n\_samples, ), fo  
r example using ravel().  
y = column\_or\_1d(y, warn=True)

## Ridge and Lasso Regression

In [42]:

```
#Ridge Regression Model
from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

In [43]:

```
plt.figure(figsize=(10,10))
```

Out[43]:

<Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

In [44]:

```
features = insurance_data.columns[0:1]
target = insurance_data.columns[-1:]
#x and y values
x = insurance_data[features].values
y = insurance_data[target].values
#split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of x\_train is (936, 1)

The dimension of x\_test is (402, 1)

In [45]:

```
#Model
lr = LinearRegression()
#fit model
lr.fit(x_train ,y_train)
#predict
#prediction = lr.predict(x_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07447061146193878

The test score for lr model is 0.10891203216512224

In [46]:

```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train,y_train)
test_score_ridge = ridgeReg.score(x_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.07446228994221393

The test score for ridge model is 0.10855133360950642

In [47]:

```
#Lasso Regression model
print("\nLasso Model:\n")
lasso = Lasso(alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls = lasso.score(x_train,y_train)
test_score_ls = lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.07446997086306062


The test score for ls model is 0.10881427793326703

In [48]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[48]:

<Axes: >



In [49]:

```
#using the Linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
#score
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.07446997086306062


0.10881427793326703

C:\Users\shaha\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:1568: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
y = column\_or\_1d(y, warn=True)



In [50]:

```
#plot size
plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label='lasso;α = grid')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation=90)
plt.legend()
plt.title("comparison plot of Ridge,Lasso and Linear Regression Model")
plt.show()
```



In [51]:

```
#using the Linear Cv model
from sklearn.linear_model import RidgeCV
#Ridge cross validation
ridge_cv = RidgeCV(alphas = [0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(x_train,y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(x_test,y_test)))
```

The train score for ridge model is 0.07446228994221393

The train score for ridge model is 0.10855133360950775

## Elastic Net

In [78]:

```
from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
regr.score(x,y)
```

```
[-5.41728687e-03 -0.00000000e+00 -0.00000000e+00 -0.00000000e+00
 -0.00000000e+00  2.81193545e-05]
0.04440232543556241
```

Out[78]:

0.6855923754478445

In [53]:

```
y_pred_elastic = regr.predict(x_train)
```

In [54]:

```
mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
print("Mean squared Error on test set",mean_squared_error)
```

Mean squared Error on test set 269213980.1466613

## Decision Tree

In [55]:

```
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

In [60]:

```
x=["age","sex","bmi","children","region"]
y=["Yes","No"]
all_inputs=df[x]
all_classes=df["smoker"]
```

In [61]:

```
(x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.30)
```

In [62]:

```
clf=DecisionTreeClassifier(random_state=0)
```

In [63]:

```
clf.fit(x_train,y_train)
```

Out[63]:

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

In [64]:

```
score=clf.score(x_test,y_test)
print(score)
```

0.6467661691542289

## Random Forest

In [65]:

```
import matplotlib.pyplot as plt,seaborn as sns
```

In [66]:

```
x=df.drop('smoker',axis=1)
y=df['smoker']
```

In [67]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7)
x_train.shape,x_test.shape
```

Out[67]:

```
((936, 6), (402, 6))
```

In [68]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[68]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [69]:

```
rf=RandomForestClassifier()
```

In [70]:

```
params={'max_depth':[2,3,5,10,20],
        'min_samples_leaf':[5,10,20,50,100,200],
        'n_estimators':[10,25,30,50,100,200]}
```

In [71]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

Out[71]:

```
► GridSearchCV
► estimator: RandomForestClassifier
  ► RandomForestClassifier
```

In [72]:

```
grid_search.best_score_
```

Out[72]:

```
0.9444444444444444
```

In [73]:

```
rf_best=grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=20, min_samples_leaf=5, n_estimators=50)
```

In [74]:

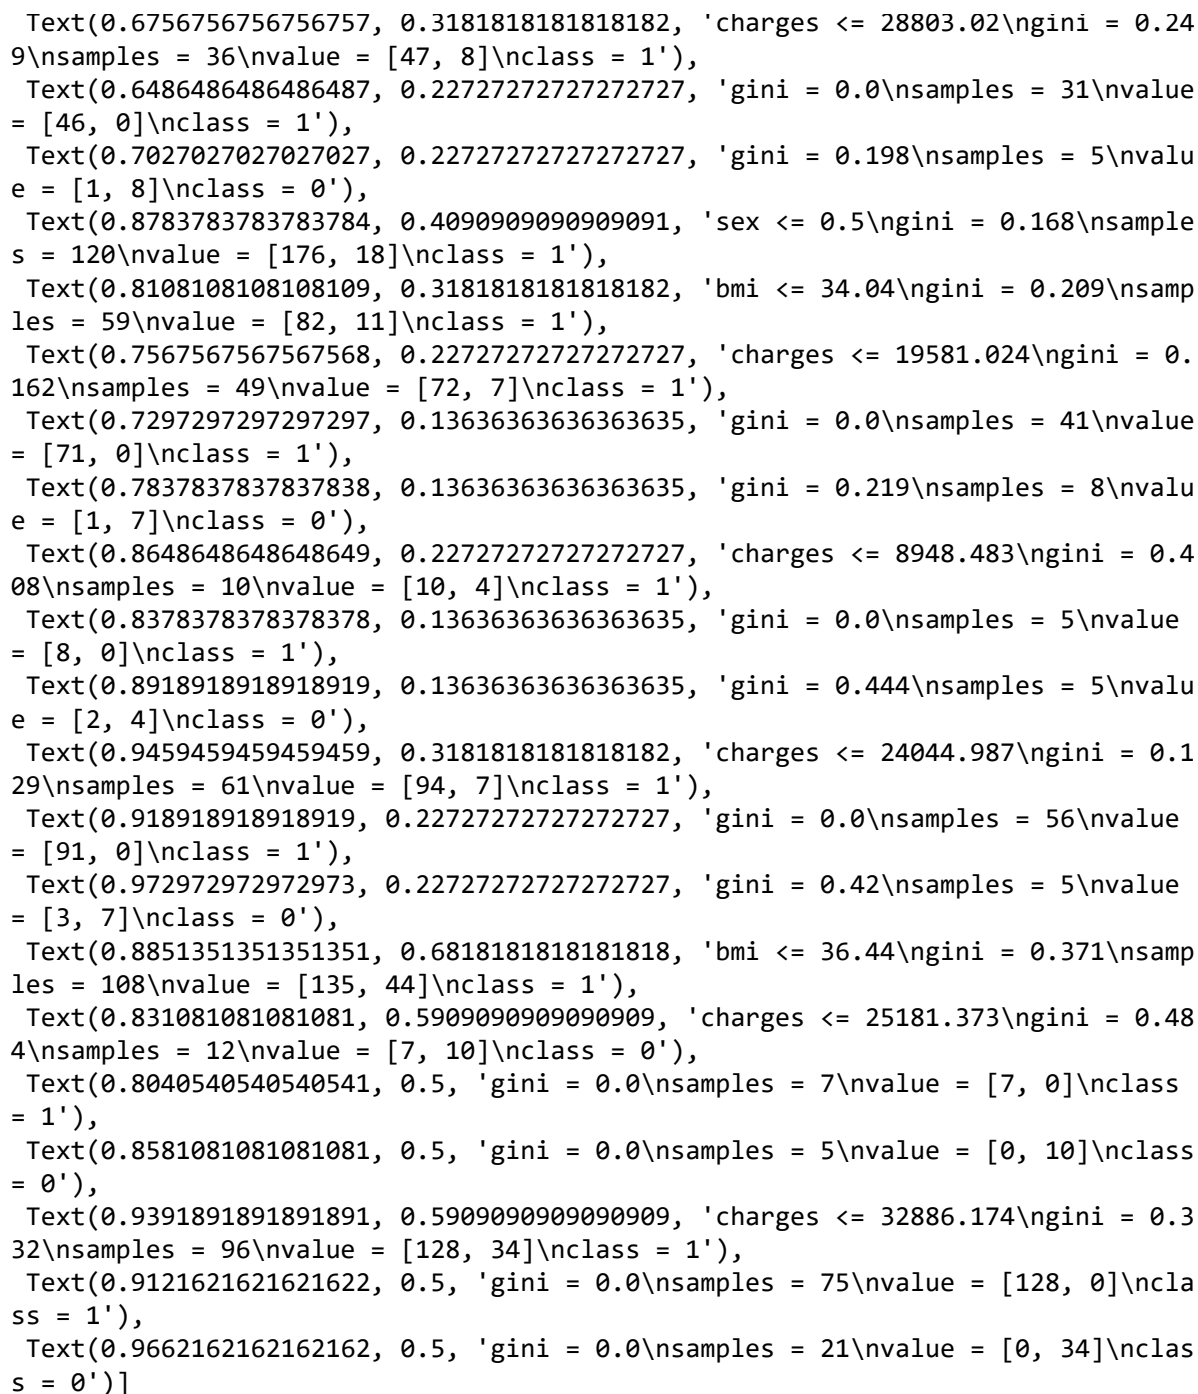
```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=["1","0"],filled=True)
```

Out[74]:

```

[Text(0.24472128378378377, 0.9545454545454546, 'bmi <= 20.517\ngini = 0.332\nsamples = 574\nvalue = [739, 197]\nclass = 1'),
Text(0.08108108108108109, 0.8636363636363636, 'bmi <= 19.998\ngini = 0.497\nsamples = 21\nvalue = [21, 18]\nclass = 1'),
Text(0.05405405405405406, 0.7727272727272727, 'children <= 1.5\ngini = 0.426\nsamples = 16\nvalue = [18, 8]\nclass = 1'),
Text(0.02702702702702703, 0.6818181818181818, 'gini = 0.165\nsamples = 7\nvalue = [10, 1]\nclass = 1'),
Text(0.08108108108108109, 0.6818181818181818, 'gini = 0.498\nsamples = 9\nvalue = [8, 7]\nclass = 1'),
Text(0.10810810810810811, 0.7727272727272727, 'gini = 0.355\nsamples = 5\nvalue = [3, 10]\nclass = 0'),
Text(0.4083614864864865, 0.8636363636363636, 'age <= 20.5\ngini = 0.319\nsamples = 553\nvalue = [718, 179]\nclass = 1'),
Text(0.16216216216216217, 0.7727272727272727, 'charges <= 14208.249\ngini = 0.438\nsamples = 68\nvalue = [75, 36]\nclass = 1'),
Text(0.13513513513513514, 0.6818181818181818, 'gini = 0.0\nsamples = 49\nvalue = [72, 0]\nclass = 1'),
Text(0.1891891891891892, 0.6818181818181818, 'sex <= 0.5\ngini = 0.142\nsamples = 19\nvalue = [3, 36]\nclass = 0'),
Text(0.16216216216216217, 0.5909090909090909, 'gini = 0.0\nsamples = 11\nvalue = [0, 26]\nclass = 0'),
Text(0.21621621621621623, 0.5909090909090909, 'gini = 0.355\nsamples = 8\nvalue = [3, 10]\nclass = 0'),
Text(0.6545608108108109, 0.7727272727272727, 'bmi <= 35.94\ngini = 0.298\nsamples = 485\nvalue = [643, 143]\nclass = 1'),
Text(0.4239864864864865, 0.6818181818181818, 'age <= 26.5\ngini = 0.273\nsamples = 377\nvalue = [508, 99]\nclass = 1'),
Text(0.2702702702702703, 0.5909090909090909, 'region <= 2.5\ngini = 0.05\nsamples = 48\nvalue = [76, 2]\nclass = 1'),
Text(0.24324324324324326, 0.5, 'gini = 0.0\nsamples = 21\nvalue = [37, 0]\nclass = 1'),
Text(0.2972972972972973, 0.5, 'charges <= 4033.209\ngini = 0.093\nsamples = 27\nvalue = [39, 2]\nclass = 1'),
Text(0.2702702702702703, 0.4090909090909091, 'gini = 0.0\nsamples = 22\nvalue = [35, 0]\nclass = 1'),
Text(0.32432432432432434, 0.4090909090909091, 'gini = 0.444\nsamples = 5\nvalue = [4, 2]\nclass = 1'),
Text(0.5777027027027027, 0.5909090909090909, 'bmi <= 28.547\ngini = 0.299\nsamples = 329\nvalue = [432, 97]\nclass = 1'),
Text(0.40540540540540543, 0.5, 'charges <= 15783.105\ngini = 0.354\nsamples = 151\nvalue = [191, 57]\nclass = 1'),
Text(0.3783783783783784, 0.4090909090909091, 'gini = 0.0\nsamples = 109\nvalue = [180, 0]\nclass = 1'),
Text(0.43243243243243246, 0.4090909090909091, 'bmi <= 22.943\ngini = 0.271\nsamples = 42\nvalue = [11, 57]\nclass = 0'),
Text(0.40540540540540543, 0.3181818181818182, 'gini = 0.49\nsamples = 5\nvalue = [3, 4]\nclass = 0'),
Text(0.4594594594594595, 0.3181818181818182, 'bmi <= 27.62\ngini = 0.228\nsamples = 37\nvalue = [8, 53]\nclass = 0'),
Text(0.43243243243243246, 0.22727272727272727, 'charges <= 22036.629\ngini = 0.287\nsamples = 29\nvalue = [8, 38]\nclass = 0'),
Text(0.3783783783783784, 0.13636363636363635, 'sex <= 0.5\ngini = 0.077\nsamples = 14\nvalue = [1, 24]\nclass = 0'),
Text(0.35135135135135137, 0.04545454545454546, 'gini = 0.0\nsamples = 7\nvalue = [0, 15]\nclass = 0'),
Text(0.40540540540540543, 0.04545454545454546, 'gini = 0.18\nsamples = 7\nvalue = [1, 9]\nclass = 0'),
Text(0.4864864864864865, 0.13636363636363635, 'age <= 55.5\ngini = 0.444\nsamples = 15\nvalue = [7, 14]\nclass = 0'),
Text(0.4594594594594595, 0.04545454545454546, 'gini = 0.355\nsamples = 9\nvalue = [3, 10]\nclass = 0'),

```





In [75]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=["Yes","No"],filled=True)
```

Out[75]:

```

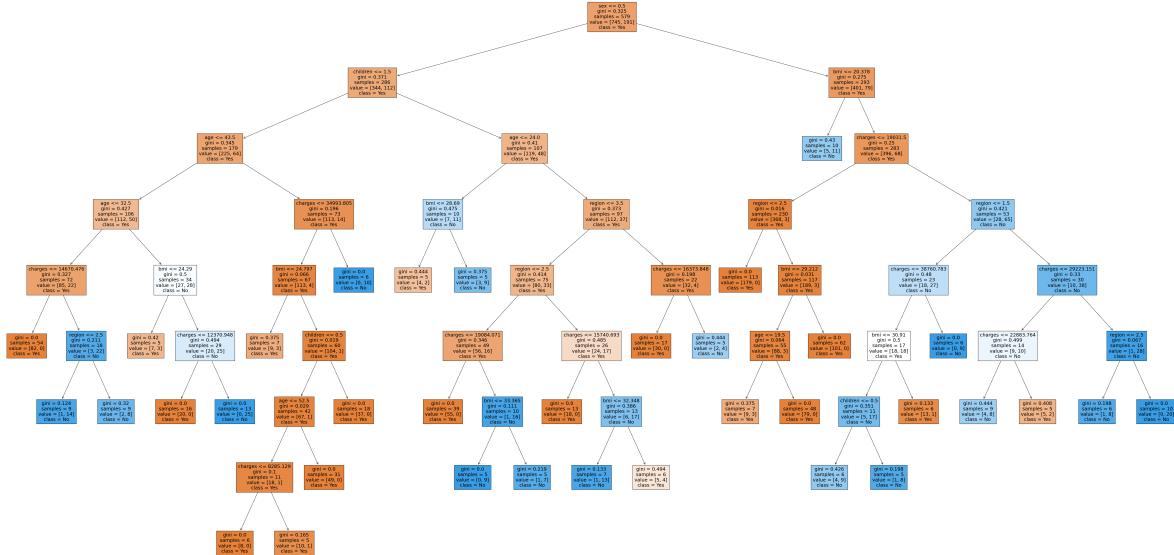
[Text(0.5171875, 0.9444444444444444, 'sex <= 0.5\ngini = 0.325\nsamples = 579\nvalue = [745, 191]\n\nclass = Yes'),
Text(0.315625, 0.8333333333333334, 'children <= 1.5\ngini = 0.371\nsamples = 286\n\nclass = Yes'),
Text(0.1875, 0.7222222222222222, 'age <= 43.5\ngini = 0.345\nsamples = 179\n\nclass = Yes'),
Text(0.1, 0.6111111111111112, 'age <= 32.5\ngini = 0.427\nsamples = 106\n\nclass = Yes'),
Text(0.05, 0.5, 'charges <= 14670.476\ngini = 0.327\nsamples = 72\n\nclass = Yes'),
Text(0.025, 0.3888888888888889, 'gini = 0.0\nsamples = 54\n\nclass = Yes'),
Text(0.075, 0.3888888888888889, 'region <= 2.5\ngini = 0.211\nsamples = 18\n\nclass = No'),
Text(0.05, 0.2777777777777778, 'gini = 0.124\nsamples = 9\n\nclass = No'),
Text(0.1, 0.2777777777777778, 'gini = 0.32\nsamples = 9\n\nclass = No'),
Text(0.15, 0.5, 'bmi <= 24.29\ngini = 0.5\nsamples = 34\n\nclass = No'),
Text(0.125, 0.3888888888888889, 'gini = 0.42\nsamples = 5\n\nclass = Yes'),
Text(0.175, 0.3888888888888889, 'charges <= 12370.948\ngini = 0.494\nsamples = 29\n\nclass = No'),
Text(0.15, 0.2777777777777778, 'gini = 0.0\nsamples = 16\n\nclass = Yes'),
Text(0.2, 0.2777777777777778, 'gini = 0.0\nsamples = 13\n\nclass = No'),
Text(0.275, 0.6111111111111112, 'charges <= 34993.805\ngini = 0.196\nsamples = 73\n\nclass = Yes'),
Text(0.25, 0.5, 'bmi <= 24.797\ngini = 0.066\nsamples = 67\n\nclass = Yes'),
Text(0.225, 0.3888888888888889, 'gini = 0.375\nsamples = 7\n\nclass = Yes'),
Text(0.275, 0.3888888888888889, 'children <= 0.5\ngini = 0.019\nsamples = 60\n\nclass = Yes'),
Text(0.25, 0.2777777777777778, 'age <= 52.5\ngini = 0.029\nsamples = 42\n\nclass = Yes'),
Text(0.225, 0.16666666666666666, 'charges <= 8285.129\ngini = 0.1\nsamples = 11\n\nclass = Yes'),
Text(0.2, 0.05555555555555555, 'gini = 0.0\nsamples = 6\n\nclass = Yes'),
Text(0.25, 0.05555555555555555, 'gini = 0.165\nsamples = 5\n\nclass = Yes'),
Text(0.275, 0.16666666666666666, 'gini = 0.0\nsamples = 31\n\nclass = Yes'),
Text(0.3, 0.2777777777777778, 'gini = 0.0\nsamples = 18\n\nclass = Yes'),
Text(0.3, 0.5, 'gini = 0.0\nsamples = 6\n\nclass = No'),
Text(0.44375, 0.7222222222222222, 'age <= 24.0\ngini = 0.41\nsamples = 107\n\nclass = Yes'),
Text(0.375, 0.6111111111111112, 'bmi <= 28.69\ngini = 0.475\nsamples = 10\n\nclass = No'),
Text(0.35, 0.5, 'gini = 0.444\nsamples = 5\n\nclass = Yes'),
Text(0.4, 0.5, 'gini = 0.375\nsamples = 5\n\nclass = No'),
Text(0.5125, 0.6111111111111112, 'region <= 3.5\ngini = 0.373\nsamples = 97\n\nclass = Yes'),
Text(0.45, 0.5, 'region <= 2.5\ngini = 0.414\nsamples = 75\n\nclass = Yes'),
Text(0.4, 0.3888888888888889, 'charges <= 19084.071\ngini = 0.346\nsamples = 49\n\nclass = Yes'),
Text(0.375, 0.2777777777777778, 'gini = 0.0\nsamples = 39\n\nclass = Yes')

```

```

ass = Yes'),
Text(0.425, 0.2777777777777778, 'bmi <= 33.365\ngini = 0.111\nsamples = 10\nvalue = [1, 16]\nclass = No'),
Text(0.4, 0.16666666666666666, 'gini = 0.0\nsamples = 5\nvalue = [0, 9]\nclass = No'),
Text(0.45, 0.16666666666666666, 'gini = 0.219\nsamples = 5\nvalue = [1, 7]\nclass = No'),
Text(0.5, 0.3888888888888889, 'charges <= 15740.693\ngini = 0.485\nsamples = 26\nvalue = [24, 17]\nclass = Yes'),
Text(0.475, 0.2777777777777778, 'gini = 0.0\nsamples = 13\nvalue = [18, 0]\nclass = Yes')

```



```

205\nvalue = [370, 00]\nclass = Yes'),
Text(0.65, 0.6111111111111112, 'region <= 2.5\ngini = 0.016\nsamples = 230\nvalue = [261, 368, 3]\nclass = Yes'),
Text(0.625, 0.5, 'gini = 0.0\nsamples = 113\nvalue = [179, 0]\nclass = Yes'),
rf_best.feature_importances_
Text(0.675, 0.5, 'bmi <= 29.212\ngini = 0.031\nsamples = 117\nvalue = [189, 3]\nclass = Yes'),

```

```

Out[76]:
Text(0.65, 0.3888888888888889, 'age <= 19.5\ngini = 0.064\nsamples = 55\nvalue = [88, 3]\nclass = Yes'),
array([0.054478, 0.010704, 0.073369, 0.014652, 0.012605,
       0.063419, 0.0977])
Text(0.625, 0.2777777777777778, 'gini = 0.375\nsamples = 7\nvalue = [9, 3]\nclass = Yes'),
Text(0.675, 0.2777777777777778, 'gini = 0.0\nsamples = 48\nvalue = [79, 0]\nclass = Yes'),

```

```

Text(0.7, 0.3888888888888889, 'gini = 0.0\nsamples = 62\nvalue = [101, 0]\nclass = Yes')
imp_df = pd.DataFrame({'varname': x_train.columns, 'Imp': rf_best.feature_importances_})
imp_df.sort_values(by="Imp", ascending=False)
Text(0.8375, 0.6111111111111112, 'region <= 1.5\ngini = 0.421\nsamples = 53\nvalue = [28, 65]\nclass = No'),

```

```

Out[77]:
Text(0.775, 0.5, 'charges <= 38760.783\ngini = 0.48\nsamples = 23\nvalue = [18, 27]\nclass = No'),
Text(0.75, 0.3888888888888889, 'bmi <= 30.91\ngini = 0.5\nsamples = 17\nvalue = [18, 18]\nclass = Yes'),
Text(0.725, 0.2777777777777778, 'children <= 0.5\ngini = 0.351\nsamples = 11\nvalue = [5, 17]\nclass = No'),
Text(0.7, 0.16666666666666666, 'gini = 0.426\nsamples = 6\nvalue = [4, 9]\nclass = No'),
Text(0.75, 0.16666666666666666, 'gini = 0.198\nsamples = 5\nvalue = [1, 8]\nclass = No'),
Text(0.775, 0.2777777777777778, 'gini = 0.133\nsamples = 6\nvalue = [13, 1]\nclass = Yes'),
Text(0.8, 0.3888888888888889, 'gini = 0.0\nsamples = 6\nvalue = [0, 9]\nclass = No'),
Text(0.9, 0.5, 'charges <= 29223.151\ngini = 0.33\nsamples = 30\nvalue = [10, 38]\nclass = No'),
Text(0.85, 0.3888888888888889, 'charges <= 22883.764\ngini = 0.499\nsamples = 14\nvalue = [9, 10]\nclass = No'),
Text(0.825, 0.2777777777777778, 'gini = 0.444\nsamples = 9\nvalue = [4, 8]\nclass = No'),

```

```
ass = No'),
In [ ]:
Text(0.875, 0.2777777777777778, 'gini = 0.408\nsamples = 5\nvalue = [5, 2]\ncl
ass = Yes'),
Text(0.95, 0.3888888888888889, 'region <= 2.5\ngini = 0.067\nsamples = 16\nval
ue = [1, 28]\nclass = No'),
In [ ]:
Text(0.925, 0.2777777777777778, 'gini = 0.198\nsamples = 6\nvalue = [1, 8]\ncl
ass = No'),
Text(0.975, 0.2777777777777778, 'gini = 0.0\nsamples = 10\nvalue = [0, 20]\ncl
ass = No')]
```

In [ ]:

In [ ]:

## Conclusion

In [ ]:

From The Given Insurance Dataset, we have formed on different models like Linear Regression, Logistic Regression, Ridge Regression, Lasso Regression, Elastic Net, Random Forest, Decision Tree. By observing the score or model prediction in this models, In Random forest model have the best score and best accuracy.

In [ ]: